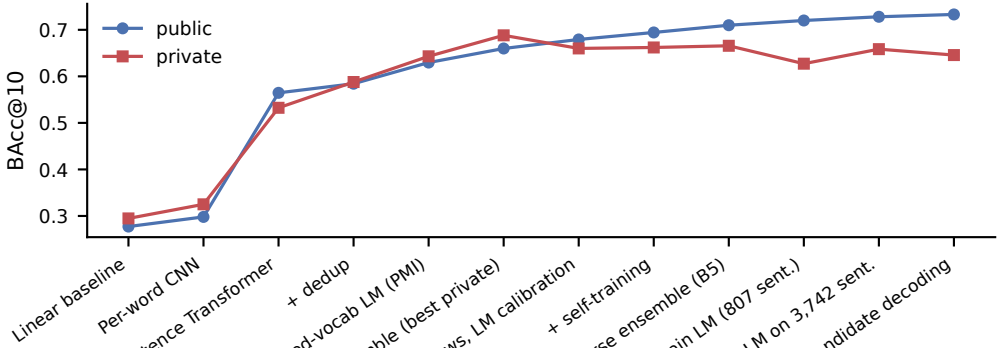
BAcc@10
0.7
0.6
0.5
0.4
0.3
public
private
Linear baseline
Per-word CNN
Sentence Transformer
+ dedup
+ closed-vocab LM (PMI)
9-model ensemble (best private)
3 s windows, LM calibration
+ self-training
+ diverse ensemble (B5)
+ in-domain LM (807 sent.)
IN1 + LM on 3,742 sent.
+ candidate decoding